

## Assignment 2: MLP and Word Vectors

Homework assignments will be done individually: each student must hand in their own answers. Use of partial or entire solutions obtained from others or online is strictly prohibited. Electronic submission on Canvas is mandatory.

### 1. Multi-Layer Perceptron (MLP) (30 pts)

- (a) Preprocess the data: tokenization, feature extraction. You can reuse the code from Assignment 1.
- (b) (20 pts) Implement the MLP class
- (c) (10 pts) Implement mini-batch GD for MLP
- (d) Run all the code to make sure your implementation works.

### 2. Word2vec - Written (25 pts)

- (a) (5 pts) Derive the gradients of the sigmoid function and show that it can be rewritten as a function of the function value (i.e., in some expressions where only  $\sigma(x)$ , but not  $x$ , is present). Assume that the input  $x$  is a scalar for this question. Recall, the sigmoid function is:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- (b) (5 pts) Assume you are given a predicted word vector  $\mathbf{v}_c$  corresponding to the center word  $c$  for skip-gram, and the word prediction is made with the softmax function:

$$\hat{y}_o = p(o|c) = \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w=1}^W \exp(\mathbf{u}_w^\top \mathbf{v}_c)}$$

where  $o$  is the expected word,  $w$  denotes the  $w$ -th word and  $\mathbf{u}_w$  ( $w = 1, \dots, W$ ) are the “output” (context) word vectors for all words in the vocabulary. The cross entropy function is defined as:

$$J_{\text{CE}}(o, \mathbf{v}_c, U) = \text{CE}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_i y_i \log(\hat{y}_i)$$

where the gold vector  $\mathbf{y}$  is a one-hot vector, the softmax prediction vector  $\hat{\mathbf{y}}$  is a probability distribution over the output space, and  $U = [u_1, u_2, \dots, u_W]$  is the matrix of all the output vectors. Assume cross entropy cost is applied to this prediction, derive the gradients with respect to  $\mathbf{v}_c$ .

- (c) (5 pts) Derive gradients for the “output” word vector  $\mathbf{u}_w$  (including  $\mathbf{u}_o$ ) in (b).
- (d) (5 pts) Repeat (b) and (c) assuming we are using the negative sampling loss for the predicted vector  $\mathbf{v}_c$ . Assume that  $K$  negative samples (words) are drawn and they are  $1, \dots, K$  respectively. For simplicity of notation, assume ( $o \notin \{1, \dots, K\}$ ). Again for a given word  $o$ , use  $\mathbf{u}_o$  to denote its output vector. The negative sampling loss function in this case is:

$$J_{\text{neg-sample}}(o, \mathbf{v}_c, U) = -\log(\sigma(\mathbf{u}_o^\top \mathbf{v}_c)) - \sum_{k=1}^K \log(\sigma(-\mathbf{u}_k^\top \mathbf{v}_c))$$

- 
- (e) (5 pts) Derive gradients for all of the word vectors for skip-gram given the previous parts and given a set of context words  $[\text{word}_{c-m}, \dots, \text{word}_c, \dots, \text{word}_{c+m}]$  where  $m$  is the context size. Denote the “input” and “output” word vectors for word  $k$  as  $\mathbf{v}_k$  and  $\mathbf{u}_k$  respectively.

*Hint:* feel free to use  $F(o, \mathbf{v}_c)$  (where  $o$  is the expected word) as a placeholder for the  $J_{\text{CE}}(o, \mathbf{v}_c \dots)$  or  $J_{\text{neg-sample}}(o, \mathbf{v}_c \dots)$  cost functions in this part – you’ll see that this is a useful abstraction for the coding part. That is, your solution may contain terms of the form  $\frac{\partial F(o, \mathbf{v}_c)}{\partial \dots}$ . Recall that for skip-gram, the cost for a context centered around  $c$  is:

$$\sum_{-m \leq j \leq m, j \neq 0} F(w_{c+j}, \mathbf{v}_c)$$

3. **Word2vec - Coding** (45 points)

- (a) (5pts) Implement Sigmoid function
- (b) (10 pts) Implement Naive Softmax loss and gradient function for word2vec models.
- (c) (10 pts) Implement Negative sampling loss function for word2vec models
- (d) (10 pts) Implement Skip-gram model in word2vec
- (e) (10 pts) Implement the k-nearest neighbors algorithm, which will be used for analysis. The algorithm receives a vector, a matrix and an integer  $k$ , and returns  $k$  indices of the matrix’s rows that are closest to the vector. Use the cosine similarity as a distance metric ([https://en.wikipedia.org/wiki/Cosine\\_similarity](https://en.wikipedia.org/wiki/Cosine_similarity)).
- (f) Load some real data and train your own word vectors. Use the training data to train word vectors. Process the dataset and use the `sgd` function and `word2vec` to generate word vectors. Visualize a few word examples. There is no additional code to write for this part.
- (g) Run the jupyter notebook code to make sure your implementation works

4. **Submission Instructions** You shall submit a zip file named `Assignment2_LastName_FirstName.zip` which contains:

- The jupyter notebook which includes all your code (and your written part), and a brief report of your knn results.
- (optional) a png (or jpg) file contains the word vector plot (`vector.png`).
- (optional) a pdf file contains all your solutions for the Written part.